

TEXT MINING

STATISTICAL MODELING OF TEXTUAL DATA

LECTURE 2

Måns Magnusson, Mattias Villani

Division of Statistics and Machine Learning
Dept. of Computer and Information Science
Linköping University

OVERVIEW

- ▶ Text classification
- ▶ Regularization
- ▶ R's tm package (demo) [TMPackageDemo.R]

SUPERVISED CLASSIFICATION

- ▶ Predict the **class label** $s \in S$ using a set of **features**.
- ▶ Feature = Explanatory variable = Predictor = Covariate
- ▶ Binary classification: $s \in \{0, 1\}$
 - ▶ Movie reviews: $S = \{\text{pos}, \text{neg}\}$
 - ▶ E-mail spam: $S = \{\text{Spam}, \text{Ham}\}$
 - ▶ Bankruptcy: $S = \{\text{Not bankrupt}, \text{Bankrupt}\}$
- ▶ Multi-class classification: $s \in \{1, 2, \dots, K\}$
 - ▶ Topic categorization of web pages:
 $S = \{\text{'News'}, \text{'Sports'}, \text{'Entertainment'}\}$
 - ▶ POS-tagging: $S = \{\text{VB}, \text{JJ}, \text{NN}, \dots, \text{DT}\}$

SUPERVISED CLASSIFICATION, CONT.

► Example data:

- Larry Wall, born in British Columbia, Canada, is the original creator of the programming language Perl. Born in 1956, Larry went to ...
- Bjarne Stroustrup is a 62-years old computer scientist ...

Person	Income	Age	Single	Payment remarks	Bankrupt
Larry	10	58	Yes	Yes	Yes
Bjarne	15	62	No	Yes	No
⋮	⋮	⋮	⋮	⋮	⋮
Guido	27	56	No	No	No

- Classification: construct prediction machine

Features \rightarrow Class label

- More generally:

Features $\rightarrow \text{Pr}(\text{Class label} | \text{Features})$


FEATURES FROM TEXT DOCUMENTS

- ▶ Any quantity computed from a document can be used as a **feature**:
 - ▶ Presence/absence of individual words
 - ▶ Number of times an individual word is used
 - ▶ Presence/absence of pairs of words
 - ▶ Presence/absence of individual bigrams
 - ▶ Lexical diversity
 - ▶ Word counts
 - ▶ Number of web links from document, possibly weighted by Page Rank.
 - ▶ etc etc

Document	has('ball')	has('EU')	has('political_arena')	wordlen	Lex. Div.	Topic
Article1	Yes	No	No	4.1	5.4	Sports
Article2	No	No	No	6.5	13.4	Sports
⋮	⋮	⋮		⋮	⋮	⋮
ArticleN	No	No	Yes	7.4	11.1	News

- ▶ Constructing clever **discriminating features** is the name of the game!

SUPERVISED LEARNING FOR CLASSIFICATION



THE BAYESIAN CLASSIFIER

- ▶ Bayesian classification

$$\operatorname{argmax}_{s \in S} p(s|\mathbf{x})$$

where $\mathbf{x} = (x_1, \dots, x_n)$ is a feature vector.

- ▶ By Bayes' theorem

$$p(s|\mathbf{x}) = \frac{p(\mathbf{x}|s)p(s)}{p(\mathbf{x})} \propto p(\mathbf{x}|s)p(s)$$

- ▶ Bayesian classification

$$\operatorname{argmax}_{s \in S} p(\mathbf{x}|s)p(s)$$

- ▶ $p(s)$ can be easily estimated from training data by relative frequencies.
- ▶ **Main problem:** Even with binary features [*has(word)*] the outcome space of $p(\mathbf{x}|s)$ is huge (=data are sparse).

NAIVE BAYES

- ▶ Naive Bayes (NB): features are assumed independent

$$p(\mathbf{x}|s) = \prod_{j=1}^n p(x_j|s)$$

- ▶ Naive Bayes solution

$$\operatorname{argmax}_{s \in S} \left[\prod_{j=1}^n p(x_j|s) \right] p(s)$$

- ▶ With binary features, $p(x_j|s)$ can be easily estimated by

$$\hat{p}(x_j|s) = \frac{C(x_j, s)}{C(s)}$$

- ▶ Example: $s = \text{news}$, $x_j = \text{has('ball')}$

$$\hat{p}(\text{has(ball)}|\text{news}) = \frac{\text{Number of news articles containing the word 'ball'}}{\text{Number of news articles}}$$

NAIVE BAYES

- ▶ **Continuous features** (e.g. lexical diversity) can be handled by:
 - ▶ Replacing continuous feature with several binary features ($1 \leq \text{lexDiv} < 2$, $2 \leq \text{lexDiv} \leq 10$ and $\text{lexDiv} > 10$)
 - ▶ Estimating $p(x_j|s)$ by a density estimator (e.g. kernel estimator)
- ▶ Finding the **most discriminatory features**. Sort from largest to smallest

$$\frac{p(x_j|s = \text{pos})}{p(x_j|s = \text{neg})} \text{ for } j = 1, \dots, n.$$

- ▶ **Problem with NB**: features are seldom independent \Rightarrow double-counting the evidence of individual features.
- ▶ **Advantages of NB**: simple and fast, yet often surprising accurate classifications.

MULTINOMIAL REGRESSION

- ▶ Logistic regression (Maximum Entropy/**MaxEnt**):

$$p(s = 1|\mathbf{x}) = \frac{\exp(\mathbf{x}'\beta)}{1 + \exp(\mathbf{x}'\beta)}$$

- ▶ Classification rule: Choose $s = 0$ if $p(s|\mathbf{x}) < 0.5$ otherwise choose $s = 1$.
- ▶ ... at least when consequences of different choices of s are the same. Loss/Utility function.
- ▶ Multinomial regression for multi-class data with K classes

$$p(s = s_j|\mathbf{x}) = \frac{\exp(\mathbf{x}'\beta_j)}{\sum_{k=1}^K \exp(\mathbf{x}'\beta_k)}$$

- ▶ Classification

$$\operatorname{argmax}_{s \in \{s_1, \dots, s_K\}} p(s|\mathbf{x})$$

- ▶ $P \times (S - 1)$ number of coefficients
- ▶ Classification with text data is like any multi-class regression problem

REGULARIZATION - VARIABLE SELECTION

- ▶ Select a subset of the covariates.
- ▶ Old school: **Forward** and **backward** selection.
- ▶ New school: **Bayesian variable selection**.
- ▶ For each β_i introduce binary indicator I_i such that

$I_i = 1$ if covariate is in the model, that is $\beta_i \neq 0$

$I_i = 0$ if covariate is in the model, that is $\beta_i = 0$

- ▶ Use Markov Chain Monte Carlo (MCMC) simulation to approximate $\Pr(I_i|Data)$ for each i .
- ▶ Example $S = \{\text{News, Sports}\}$. $\Pr(\text{News}|x)$.

	has('ball')	has('EU')	has('political_arena')	wordlen	Lex. Div.
$\Pr(I_i Data)$	0.2	0.90	0.99	0.01	0.85

REGULARIZATION - SHRINKAGE

- ▶ Keep all covariates, but **shrink** their β -coefficient to zero.
- ▶ **Penalized likelihood**

$$L_{Ridge}(\beta) = LogLik(\beta) - \lambda \beta' \beta$$

where λ is the **penalty parameter**.

- ▶ Maximize $L_{Ridge}(\beta)$ with respect to β . Trade-off of fit ($LogLik(\beta)$) against complexity penalty $\beta' \beta$.
- ▶ **Ridge regression** if regression is linear.
- ▶ The penalty can be motivated as a **Bayesian prior** $\beta_i \stackrel{iid}{\sim} N(0, \lambda^{-1})$.
- ▶ λ can be estimated by cross-validation or Bayesian methods.

LASSO - SHRINKAGE AND VARIABLE SELECTION

- ▶ Replace Ridge penalty

$$L_{Ridge}(\beta) = \text{LogLik}(\beta) - \lambda \sum_{j=1}^n \beta_j^2$$

by

$$L_{Lasso}(\beta) = \text{LogLik}(\beta) - \lambda \sum_{j=1}^n |\beta_j|$$


- ▶ The β that maximizes $L_{Lasso}(\beta)$ is called the **Lasso estimator**.
- ▶ Some parameters are shrunk exactly to zero \Rightarrow Lasso does **both shrinkage AND variable selection**.
- ▶ Lasso penalty is equivalent to a double exponential prior

$$p(\beta_i) = \frac{\lambda}{2} \exp(\lambda |\beta_i - 0|)$$

SUPPORT VECTOR MACHINES

- ▶ One of the best off-the-shelf classifiers around.
- ▶ Finds the line in covariate space that maximally separates the two classes.
- ▶ When the points are not linearly separable: add a slack-variable $\xi_i > 0$ for each observation. Allow misclassification, but make it costly.
- ▶ Non-linear separating curves can be obtained by basis expansion (think about adding x^2 , x^3 and so on)
- ▶ The kernel trick makes it possible to handle many covariates.
- ▶ **Drawback:** not so easily extended to multi-class.
- ▶ `svm` function in R-package `e1071` [or `nltk.classify.svm`]


LINEAR SVMs



REGRESSION TREES AND RANDOM FOREST

- ▶ Binary partitioning tree.
- ▶ At each internal node decide:
 - ▶ Which covariate to split on
 - ▶ Where to split the covariate ($X_j < c$. Trivial for binary covariates)
- ▶ The optimal splitting variables and split-points are chosen to minimize the mis-classification rate (or other similar measures).
- ▶ **Random forest (RF)** predicts using an average of many small trees.
- ▶ Each tree in RF is grown on a random subset of variables. Makes it possible to handle **many covariates**. Parallel.
- ▶ Advantage of RF: better predictions than trees.
- ▶ RF harder to interpret, but provide variable importance scores.
- ▶ **R packages:** `tree` and `rpart` (trees), `randomForest` (RF).

REGRESSION TREES



EVALUATING A CLASSIFIER: ACCURACY AND ERROR

- **Confusion matrix:**

		Truth	
		Spam	Not Spam
Decision	Spam	tp	fp
	Not Spam	fn	tn

- tp = true positive, fp = false positive
- fn = false negative, tn = true negative
- **Accuracy** is the proportion of correctly classified items

$$\text{Accuracy} = \frac{tp + tn}{tp + tn + fn + fp}$$

- **Error** is the proportion of wrongly classified items

$$\text{Error} = 1 - \text{Accuracy}$$

ACCURACY CAN BE MISLEADING

- Accuracy is problematic when t_n is large. High accuracy can then be obtained by not acting at all!

		Truth	
		Spam	Not Spam
Choice	Spam	0	0
	Not Spam	100	900

EVALUATING A CLASSIFIER: THE F-MEASURE

- **Confusion matrix:**

		Truth	
		Spam	Good
Choice	Spam	tp	fp
	Good	fn	tn

- **Precision** = proportion of selected items that the system got right

$$\text{Precision} = \frac{tp}{tp+fp}$$

- **Recall** = proportion of spam that the system classified as spam

$$\text{Recall} = \frac{tp}{tp+fn}$$

- **F-measure** is a trade-off between Precision and Recall (harmonic mean):

$$F = \frac{1}{\alpha \frac{1}{\text{Precision}} + (1 - \alpha) \frac{1}{\text{Recall}}}$$